

One output-bias gradient step (eta=10); no base-weight fine-tuning

